

Data Visualization CHIS Exercise

John Campi

November 14, 2018

Contents

1	Case Study	1
1.1	CHIS - Descriptive Statistics	1
1.2	Exploring Data	1
1.3	Unusual Values	5
1.4	Default Binwidths	5
1.5	Data Cleaning	5
1.6	Multiple Histograms	6
1.7	Alternatives	7
1.8	Do Things Manually	13
1.9	Marimekko/Mosaic Plot	15
1.10	Adding statistics	17
1.11	Adding text	19
1.12	Generalizations	22

1 Case Study

In this case study, we'll explore the large, publicly available California Health Interview Survey dataset from 2009. We'll go step-by-step through the development of a new plotting method - a mosaic plot - combining statistics and flexible visuals. At the end, we'll generalize our new plotting method to use on a variety of datasets we've seen throughout the first two courses.

1.1 CHIS - Descriptive Statistics

See video. Original dataset located at (<http://healthpolicy.ucla.edu>) Modified dataset (adult.csv) located in Data directory.

1.2 Exploring Data

In this chapter we're going to continuously build on our plotting functions and understanding to produce a mosaic plot (aka Marimekko plot). This is a visual representation of a contingency table, comparing two categorical variables. Essentially, our question is which groups are over or under represented in our dataset. To visualize this we'll color groups according to their Pearson residuals from a chi-squared test. At the end of it all we'll wrap up our script into a flexible function so that we can look at other variables.

We'll familiarize ourselves with a small number of variables from the 2009 CHIS adult-response dataset (as opposed to children). We have selected the following variables to explore:

- RBMI: BMI Category description
- BMI_P: BMI value
- RACEHPR2: race
- SRSEX: sex
- SRAGE_P: age
- MARIT2: Marital status
- AB1: General Health Condition
- ASTCUR: Current Asthma Status

- AB51: Type I or Type II Diabetes
- POVLL: Poverty level

We'll filter our dataset to plot a more reliable subset (we'll still retain over 95% of the data).

Before we get into mosaic plots it's worthwhile exploring the dataset using simple distribution plots - i.e. histograms.

ggplot2 is already loaded and the dataset, named `adult`, is already available in the workspace.

Instructions

- Use the typical commands for exploring the structure of `adult` to get familiar with the variables: `summary()` and `str()`.
- As a first exploration of the data, plot two histograms using `ggplot2` syntax: one for age (`SRAGE_P`) and one for BMI (`BMI_P`). The goal is to explore the dataset and get familiar with the distributions here. Feel free to explore different bin widths. We'll ask some questions about these in the next exercises.
- Next plot a binned-distribution of age, filling each bar according to the BMI categorization. Inside `geom_histogram()`, set `binwidth = 1`. You'll want to use `fill = factor(RBMI)` since `RBMI` is a categorical variable.

```
# Read 2009 CHIS adult-response dataset.
path_in <- file.path("Data", "adult.csv")
adult <- read.csv(path_in, stringsAsFactors=FALSE, na.strings = "")

# Explore the dataset with summary and str
summary(adult)
```

```
##      RBMI      BMI_P      RACEHPR2      SRSEX
## Min.   :1.000   Min.   :12.65   Min.   :1.000   Min.   :1.000
## 1st Qu.:2.000   1st Qu.:22.77   1st Qu.:5.000   1st Qu.:1.000
## Median :3.000   Median :25.72   Median :6.000   Median :2.000
## Mean   :2.748   Mean   :26.64   Mean   :5.088   Mean   :1.591
## 3rd Qu.:3.000   3rd Qu.:29.32   3rd Qu.:6.000   3rd Qu.:2.000
## Max.   :4.000   Max.   :93.72   Max.   :6.000   Max.   :2.000
##      SRAGE_P      MARIT2      AB1      ASTCUR
## Min.   :18.00   Min.   :1.000   Min.   :1.000   Min.   :1.000
## 1st Qu.:44.00   1st Qu.:1.000   1st Qu.:2.000   1st Qu.:2.000
## Median :57.00   Median :1.000   Median :2.000   Median :2.000
## Mean   :56.14   Mean   :2.043   Mean   :2.525   Mean   :1.915
## 3rd Qu.:69.00   3rd Qu.:3.000   3rd Qu.:3.000   3rd Qu.:2.000
## Max.   :85.00   Max.   :4.000   Max.   :5.000   Max.   :2.000
##      AB51      POVLL
## Min.   : -1.0000   Min.   :1.000
## 1st Qu.: -1.0000   1st Qu.:2.000
## Median : -1.0000   Median :4.000
## Mean   : -0.7108   Mean   :3.196
## 3rd Qu.: -1.0000   3rd Qu.:4.000
## Max.   : 3.0000   Max.   :4.000
```

```
str(adult)
```


```
## 'data.frame':   44346 obs. of  10 variables:
## $ RBMI      : int  3 3 3 2 3 4 3 2 3 3 ...
## $ BMI_P     : num  28.9 26.1 25.1 25 25.1 ...
## $ RACEHPR2  : int  6 6 6 6 6 6 6 6 6 ...
## $ SRSEX     : int  1 2 1 1 1 2 1 2 1 2 ...
```

```
## $ SRAGE_P : int 32 80 71 39 75 53 42 33 67 52 ...
## $ MARIT2  : int 1 3 1 4 1 1 1 1 3 3 ...
## $ AB1      : int 1 1 2 1 2 3 2 2 1 5 ...
## $ ASTCUR   : int 2 2 1 2 2 1 2 2 2 2 ...
## $ AB51     : int -1 -1 -1 -1 -1 -1 -1 -1 -1 -1 ...
## $ POVLL    : int 4 4 4 4 4 4 4 3 4 4 ...
```

```
# Age histogram
```

```
ggplot(adult, aes(x = SRAGE_P)) +
  geom_histogram()
```


```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```




```
# BMI value histogram
```

```
ggplot(adult, aes(x = BMI_P)) +
  geom_histogram()
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



```
# Age colored by BMI, binwidth = 1  
ggplot(adult, aes(x = SRAGE_P, fill=factor(RBMI))) +  
  geom_histogram(binwidth=1)
```



1.3 Unusual Values

In the previous exercise you used histograms to explore the age and BMI distributions and their relationships to each other in the CHIS dataset. What unusual phenomenon stood out?

If you want to experiment some more with the data, go ahead - it's available as `adult` in your workspace!

Instructions

Possible Answers (1) BMI is symmetrical around 25, the border between healthy and overweight. (2) There is an unexpectedly large number of very old people. (3) The proportion of each BMI category is consistent throughout ages. (4) The dataset indicates that California is a very young population.

Answer: (2) There is an unexpectedly large number of very old people.

1.4 Default Binwidths

If you don't specify the `binwidth` argument inside `geom_histogram()` you can tell from the message that 30 bins are used by default. This will then specify the binwidth that is used. What is this binwidth for the age variable, `SRAGE_P`, of the `adult` dataset?

Instructions

Possible Answers (1) 1.00 (2) 1.66 (3) 2.00 (4) 2.23

Answer: $(85 - 18)/30 = 2.23$

1.5 Data Cleaning

Now that we have an idea about our data, let's clean it up.

Instructions

- You should have noticed in the age distribution that there is an unusual spike of individuals at 85, which seems like an artifact of data collection and storage. Solve this by only keeping observations for which `adult$SRAGE_P` is smaller than or equal to 84.
- There is a long positive tail on the BMIs that we'd like to remove. Only keep observations for which `adult$BMI_P` is larger than or equal to 16 and `adult$BMI_P` is strictly smaller than 52.
- We'll focus on the relationship between the BMI score (& category), age and race. To make plotting easier later on, we'll change the labels in the dataset. Define `adult$RACEHPR2` as a factor with labels `c("Latino", "Asian", "African American", "White")` using the labels `c("Under-weight", "Normal-weight", "Over-weight", "Obese")`.

```
# Keep adults younger than or equal to 84
adult <- adult[adult$SRAGE_P <= 84, ]

# Keep adults with BMI at least 16 and less than 52
adult <- adult[adult$BMI_P >= 16 & adult$BMI_P < 52, ]

# Relabel the race variable
adult$RACEHPR2 <- factor(adult$RACEHPR2, labels = c("Latino", "Asian", "African American", "White"))

# Relabel the BMI categories variable
adult$RBMI <- factor(adult$RBMI, labels = c("Under-weight", "Normal-weight", "Over-weight", "Obese"))
```

1.6 Multiple Histograms

When we introduced histograms we focused on univariate data, which is exactly what we've been doing here. However, when we want to explore distributions further there is much more we can do. For example, there are density plots, which you'll explore in the next course. For now, we'll look deeper at frequency histograms and begin developing our mosaic plots.

The adult dataset, which is cleaned up by now, is available in the workspace for you.

Two layers have been pre-defined for you: `BMI_fill` is a scale layer which we can add to a `ggplot()` command using `+`: `ggplot(...) + BMI_fill`. `fix_strips` is a `theme()` layer to make nice facet titles.

Instructions

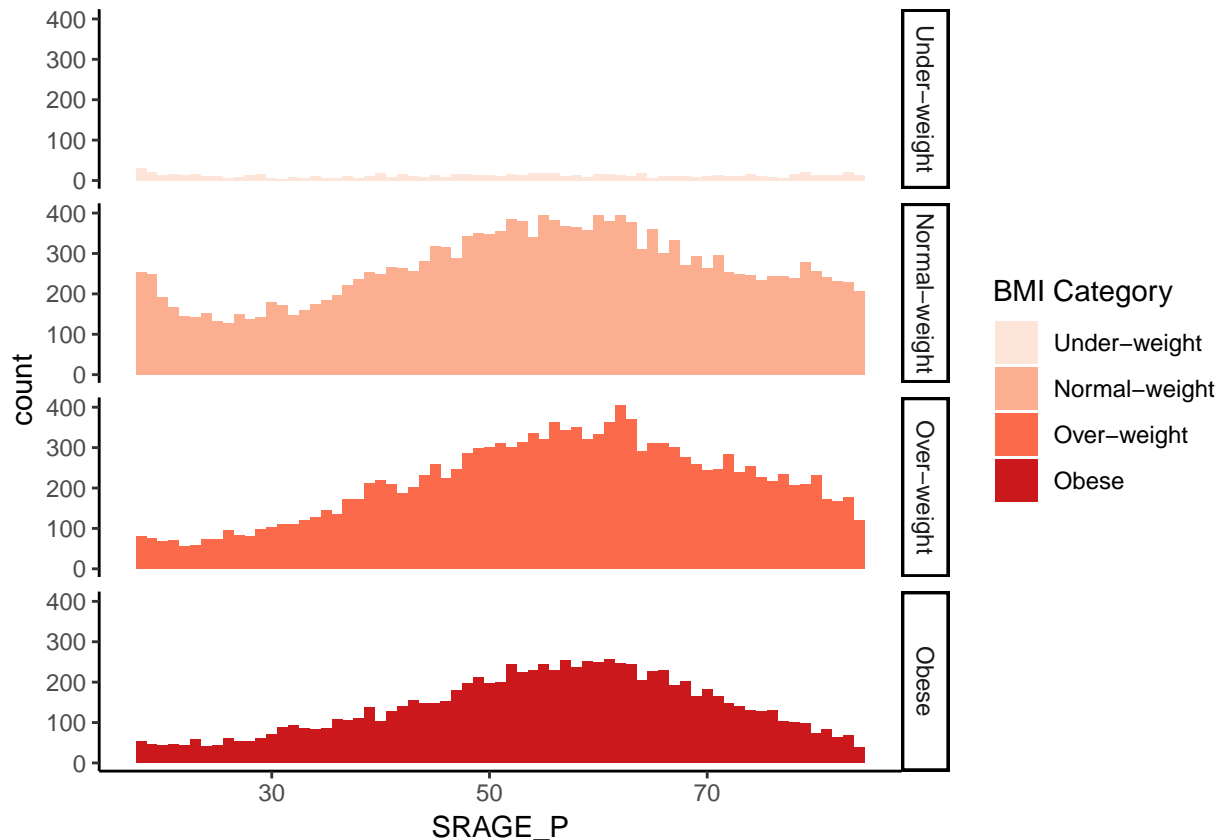
- The histogram from the first exercise of age colored by BMI has been provided. The predefined `theme()`, `fix_strips` (see above), has been added to the histogram. Add `BMI_fill` to this plot using the `+` operator as well.
- In addition, add the following elements to create a pretty & insightful plot:
- Use `facet_grid()` to facet the rows according to `RBMI` (Remember formula notation `ROWS ~ COL` and use `.` as a place-holder when not faceting in that direction).
- Add the classic theme using `theme_classic()`.

```
# The dataset adult is available

# The color scale used in the plot
BMI_fill <- scale_fill_brewer("BMI Category", palette = "Reds")

# Theme to fix category display in faceted plot
fix_strips <- theme(strip.text.y = element_text(angle = 0,
                                                  hjust = 0,
                                                  vjust = 0.1,
                                                  size = 14),
                  strip.background = element_blank(),
                  legend.position = "none")
```

```
# Histogram, add BMI_fill and customizations
ggplot(adult, aes (x = SRAGE_P, fill= factor(RBMI))) +
  geom_histogram(binwidth = 1) +
  fix_strips +
  BMI_fill +
  facet_grid(RBMI ~ .) +
  theme_classic()
```



1.7 Alternatives

In the previous exercise we looked at different ways of showing the absolute count of multiple histograms. This is fine, but density would be a more useful measure if we wanted to see how the frequency of one variable changes across another. However, there are some difficulties here, so let's take a closer look at different plots.


The clean adult dataset is available, as is the BMI_fill color palette. The first plot simply shows a histogram of counts, without facets, without modified themes. It's denoted Plot 1.

Instructions


- Plot 2 - Copy, paste and adapt the code for plot 1 so that it shows density. Do this by adding `aes(y = ..density..)` inside the `geom_histogram()` function. This plot looks really strange, because we get the density within each BMI category, not within each age group!
- Plot 3 - starting from plot 1, create a faceted histogram. Use `facet_grid()` with the formula: `RBMI ~ ..`
- Plot 4 - starting from plot 2, create a faceted histogram showing density. Use `facet_grid()` with the formula `RBMI ~ ..`. Plots 3 and 4 can be useful if we are interested in the frequency distribution within each BMI category.

- Plot 5 - Change the second plot to have position = “fill”. This is not an accurate representation, as density calculates the proportion across category, and not across bin.
- Plot 6 - To get an accurate visualization, change Plot 5, but this time, instead of `..density..`, set the y aesthetic to `..count../sum(..count..)`.


```
# Plot 1 - Count histogram
ggplot(adult, aes (x=SRAGE_P, fill=factor(RBMI))) +
  geom_histogram(binwidth=1) +
  BMI_fill
```




```
# Plot 2 - Density histogram
ggplot(adult, aes (x=SRAGE_P, fill=factor(RBMI))) +
  geom_histogram(aes(y=..density..), binwidth=1) +
  BMI_fill
```


```
# Plot 3 - Faceted count histogram
ggplot(adult, aes (x=SRAGE_P, fill=factor(RBMI))) +
  geom_histogram(binwidth=1) +
  BMI_fill +
  facet_grid(RBMI ~ .)
```




```
# Plot 4 - Faceted density histogram
ggplot(adult, aes (x=SRAGE_P, fill=factor(RBMI))) +
  geom_histogram(aes(y=..density..), binwidth=1) +
  BMI_fill +
  facet_grid(RBMI ~ .)
```



```
# Plot 5 - Density histogram with position = "fill"
ggplot(adult, aes (x=SRAGE_P, fill=factor(RBMI))) +
  geom_histogram(aes(y=..density..), binwidth=1, position="fill") +
  BMI_fill
```



```
# Plot 6 - The accurate histogram
ggplot(adult, aes (x = SRAGE_P, fill= factor(RBMI))) +
  geom_histogram(aes(y = ..count../sum(..count..)), binwidth = 1, position="fill") +
  BMI_fill
```



1.8 Do Things Manually

In the previous exercise we looked at how to produce a frequency histogram when we have many sub-categories. The problem here is that this can't be faceted because the calculations occur on the fly inside ggplot2.

To overcome this we're going to calculate the proportions outside ggplot2. This is the beginning of our flexible script for a mosaic plot.

The dataset `adult` and the `BMI_fill` object from the previous exercise have been carried over for you. Code that tries to make the accurate frequency histogram faceted is available. You should understand these commands by now.

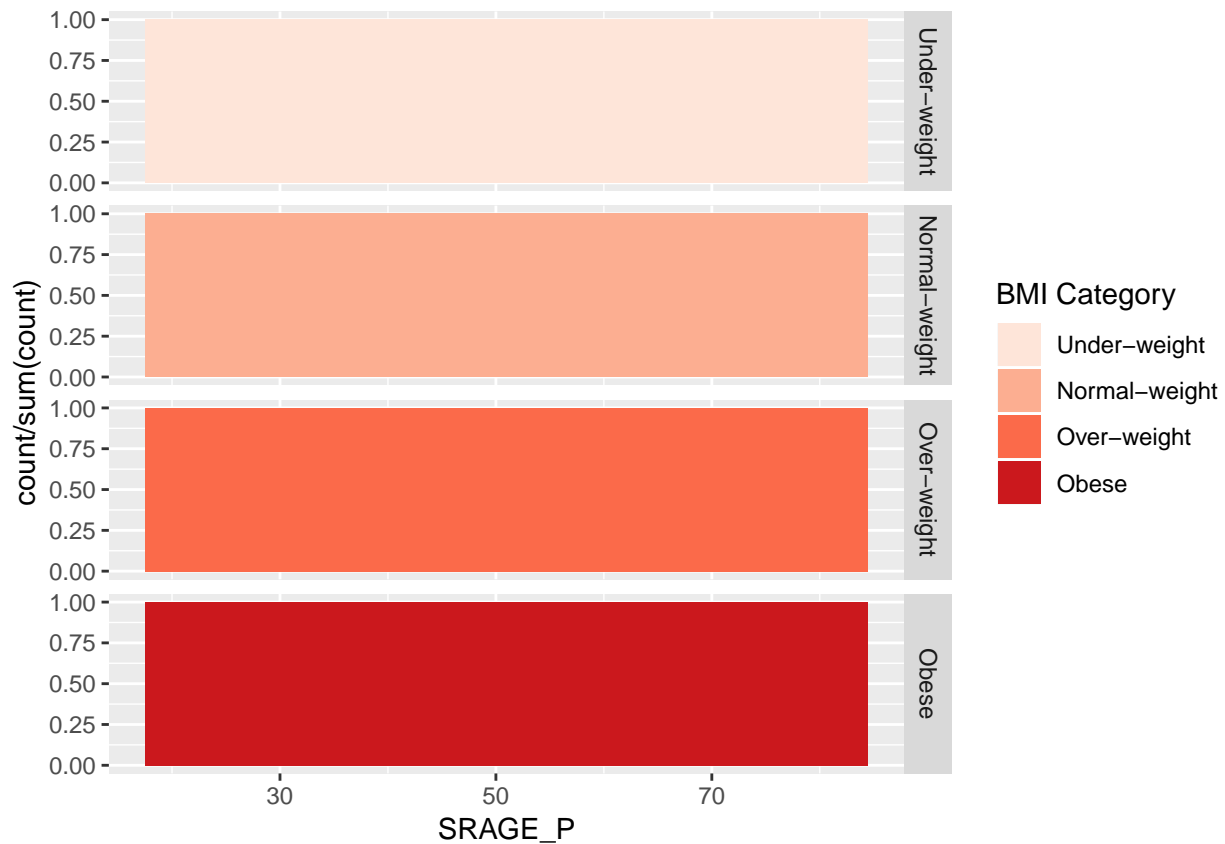
Instructions

- Use `adult` and `BMI_fill` as arguments in `table()` to create a contingency table of the two variables. Save this as `DF`.
- Use `apply()` To get the frequency of each group. The first argument is `DF`, the second argument 2, because you want to do calculations on each column. The third argument should be `function(x) x/sum(x)`. Store the result as `DF_freq`.
- Load the `reshape2` package and use the `melt()` function on `DF_freq`. Store the result as `DF_melted`. Examine the structure of `DF_freq` and `DF_melted` if you are not familiar with this operation.

Note: Here we use `reshape2` instead of the more current `tidyr` because `reshape2::melt()` allows us to work directly on a table. `tidyr::gather()` requires a data frame.

- Use `names()` to rename the variables in `DF_melted` to be `c("FILL", "X", "value")`, with the prospect of making this a generalized function later on.
- The plotting call at the end uses `DF_melted`. Add code to make it faceted. Use the formula `FILL ~ ..`. Note that we use `geom_col()` now, this is just a short-cut to `geom_bar(stat = "identity")`.

```
# An attempt to facet the accurate frequency histogram from before (failed)
ggplot(adult, aes (x=SRAGE_P, fill=factor(RBMI))) +
  geom_histogram(aes(y = ..count../sum(..count..)), binwidth=1, position="fill") +
  BMI_fill +
  facet_grid(RBMI ~ .)
```




```
# Create DF with table()
DF <- table(adult$RBMI, adult$SRAGE_P)

# Use apply on DF to get frequency of each group
DF_freq <- apply(DF, 2, function(x) x/sum(x))

# Load reshape2 and use melt on DF to create DF_melted
library(reshape2)
DF_melted <- melt(DF_freq)

# Change names of DF_melted
names(DF_melted) <- c("FILL", "X", "value")

# Add code to make this a faceted plot
ggplot(DF_melted, aes(x=X, y=value, fill=FILL)) +
  geom_col(position="stack") +
  BMI_fill +
  facet_grid(FILL ~ .) # Facets
```



1.9 Marimekko/Mosaic Plot

In the previous exercise we looked at different ways of showing the frequency distribution within each BMI category. This is all well and good, but the absolute number of each age group also has an influence on if we will consider something as over-represented or not. Here, we will proceed to change the widths of the bars to show us something about the n in each group.

This will get a bit more involved, because the aim is not to draw bars, but rather rectangles, for which we can control the widths. You may have already realized that bars are simply rectangles, but we don't have easy access to the `xmin` and `xmax` aesthetics, but in `geom_rect()` we do! Likewise, we also have access to `ymin` and `ymax`. So we're going to draw a box for every one of our 268 distinct groups of BMI category and age.

The clean adult dataset, as well as `BMI_fill`, are already available. Instead of running `apply()` like in the previous exercise, the contingency table has already been transformed to a data frame using `as.data.frame.matrix()`.

Instructions

- To build the rectangle plot, we'll add several variables to DF:
- `groupSum`, containing the sum of each row in the DF. Use `rowSums()` to calculate this. `groupSum` represents the total number of individuals in each age group.
- `xmax`: the `xmax` value for each rectangle, calculated as `cumsum(DF$groupSum)`
- `xmin`: the `xmin` value for each rectangle, calculated by subtracting the `groupSum` column from the `xmax` column.
- The names of the x axis groups are stored in the row names, which is pretty bad style, so make a new variable, `X`, that stores the values of `row.names()` for DF.
- Now we are ready to melt the dataset. Load `reshape2` and use `melt()` on DF. Specify the `id.vars` variables as `c("X", "xmin", "xmax")` and the `variable.name` argument as `"FILL"`. Store the result as `DF_melted`.

- Have a look at the dplyr call that calculates the ymax and ymin columns of DF_melted. It first groups by X and then calculates cumulative proportions. The result is stored as DF_melted again.
- If all goes well you should see the plot in the viewer when you execute the plotting function at the bottom of the script.

```
# The initial contingency table
DF <- as.data.frame.matrix(table(adult$SRAGE_P, adult$RBMI))

# Create groupSum, xmax and xmin columns
DF$groupSum <- rowSums(DF)
DF$xmax <- cumsum(DF$groupSum)
DF$xmin <- DF$xmax - DF$groupSum

# The groupSum column needs to be removed; don't remove this line
DF$groupSum <- NULL

# Copy row names to variable X
DF$X <- row.names(DF)

# Melt the dataset
library(reshape2)
DF_melted <- melt(DF, id.vars=c("X", "xmin", "xmax"), variable.name="FILL")

# dplyr call to calculate ymin and ymax - don't change
library(dplyr)


##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

DF_melted <- DF_melted %>%
  group_by(X) %>%
  mutate(ymax = cumsum(value/sum(value)),
         ymin = ymax - value/sum(value))

# Plot rectangles - don't change
library(ggthemes)
ggplot(DF_melted, aes(ymin = ymin,
                     ymax = ymax,
                     xmin = xmin,
                     xmax = xmax,
                     fill = FILL)) +
  geom_rect(colour = "white") +
  scale_x_continuous(expand = c(0,0)) +
  scale_y_continuous(expand = c(0,0)) +
  BMI_fill +
  theme_tufte()
```

1.10 Adding statistics

In the previous exercise we generated a plot where each individual bar was plotted separately using rectangles (shown in the viewer). This means we have access to each piece and we can apply different fill parameters.

So let's make some new parameters. To get the Pearson residuals, we'll use the `chisq.test()` function.

The data frames `adult` and `DF_melted`, as well as the object `BMI_fill` that you created throughout this chapter, are all still available. The `reshape2` package is already loaded.

Instructions

- Use the `adult$RBMi` (corresponding to `FILL`) and `adult$SRAGE_P` (corresponding to `X`) columns inside the `table()` function that's inside the `chisq.test()` function. Store the result as `results`.
- The residuals can be accessed through `results$residuals`. Apply the `melt()` function on them with no further arguments. Store the resulting data frame as `resid`.
- Change the names of `resid` to `c("FILL", "X", "residual")`. This is so that we have a consistent naming convention similar to how we called our variables in the previous exercises.
- The data frame from the previous exercise, `DF_melted` is already available. Use the `merge()` function to bring the two data frames together. Store the result as `DF_all`.
- Adapt the code in the `ggplot` command to use `DF_all` instead of `DF_melted`. Also, map `residual` onto `fill` instead of `FILL`.

```
# Perform chi.sq test (RBMi and SRAGE_P)
results <- chisq.test(table(adult$RBMi, adult$SRAGE_P))

# Melt results$residuals and store as resid
resid <- melt(results$residuals)
```

```

# Change names of resid
names(resid) <- c("FILL", "X", "residual")


# merge the two datasets:
DF_all <- merge(DF_melted, resid)

# Update plot command
library(ggthemes)
p <- ggplot(DF_all, aes(ymin = ymin,
                        ymax = ymax,
                        xmin = xmin,
                        xmax = xmax,
                        fill = residual)) +

  geom_rect() +
  scale_fill_gradient2() +
  scale_x_continuous(expand = c(0,0)) +
  scale_y_continuous(expand = c(0,0)) +
  theme_tufte()

# Resulting plot.
p

```



1.11 Adding text

Since we're not coloring according to BMI, we have to add group (and x axis) labels manually. Our goal is the plot in the viewer.

For this we'll use the label aesthetic inside `geom_text()`. The actual labels are found in the `FILL` (BMI category) and `X` (age) columns in the `DF_all` data frame. (Additional attributes have been set inside `geom_text()` in the exercise for you).

The labels will be added to the right (BMI category) and top (age) inner edges of the plot. (We could have also added margin text, but that is a more advanced topic that we'll encounter in the third course. This will be a suitable solution for the moment.)

The first two commands show how we got the the four positions for the y axis labels. First, we got the position of the maximum `xmax` values, i.e. at the very right end, stored as `index`. We want to calculate the half difference between each pair of `ymax` and `ymin` (e.g. $(ymax - ymin)/2$) at these index positions, then add this value to the `ymin` value. These positions are stored in the variable `yposn`.


We'll begin with the plot thus far, stored as object `p`. In the sample code, `%+%` `DF_all` refreshes the plot's dataset with the extra columns.

Instructions

- Plot 1: In the `geom_text()` function, define the x, y and label aesthetics.
- Set `x` to `max(xmax)`, so the labels are on the right side of the plot.
- Set the position of `y` to `yposn`.
- Set the label text to `FILL`.
- Plot 2: The same thing for the x axis label positions. You don't need to find an index here, since you can use the same y position for all these labels: 1.
- Calculate the half difference between each pair of `xmax` and `xmin` then add this value to `xmin`.
- Complete the plot command by adding the labels in the `xposn` to our plot, the label this time will be `X`, which in this case is the age.


This plot isn't perfect, but it does a pretty good job for an exploratory plot.

```
# Plot so far
p
```




```
# Position for labels on y axis (don't change)
index <- DF_all$xmax == max(DF_all$xmax)
DF_all$yposn <- DF_all$ymin[index] + (DF_all$ymax[index] - DF_all$ymin[index])/2

# Plot 1: geom_text for BMI (i.e. the fill axis)
p1 <- p %>% DF_all +
  geom_text(aes(x = max(xmax), y = yposn, label = FILL),
    size = 3, hjust = 1,
    show.legend = FALSE)
p1
```



```
# Plot 2: Position for labels on x axis
DF_all$xposn <- DF_all$xmin + (DF_all$xmax - DF_all$xmin)/2

# geom_text for ages (i.e. the x axis)
p1 %>% DF_all +
  geom_text(aes(x = xposn, label = X),
            y = 1, angle = 90,
            size = 3, hjust = 1,
            show.legend = FALSE)
```



1.12 Generalizations

Now that you’ve done all the steps necessary to make our mosaic plot, you can wrap all the steps into a single function that we can use to examine any two variables of interest in our data frame (or in any other data frame for that matter). For example, we can use it to examine the Vocab data frame we saw earlier in this course.

You’ve seen all the code in our function, so there shouldn’t be anything surprising there. Notice that the function takes multiple arguments, such as the data frame of interest and the variables that you want to create the mosaic plot for. None of the arguments have default values, so you’ll have to specify all three if you want the mosaicGG() function to work.

Start by going through the code and see if you understand the function’s implementation.

Instructions

- Print mosaicGG and read its contents.
- Calling mosaicGG(adult, “SRAGE_P”, “RBMI”) will result in the plot you’ve been working on so far. Try this out. This gives you a mosaic plot where BMI is described by age.
- Test out another combination of variables in the adult data frame: Poverty (POVLL) described by Age (SRAGE_P).
- Try the function on other datasets we’ve worked with throughout this course:
- mtcars dataset: am described by cyl
- Vocab dataset: vocabulary described by education.

```
# Load all packages
library(ggplot2)
library(reshape2)
```

```

library(dplyr)
library(ggthemes)

# Script generalized into a function
mosaicGG <- function(data, X, FILL) {
  # Proportions in raw data
  DF <- as.data.frame.matrix(table(data[[X]], data[[FILL]]))
  DF$groupSum <- rowSums(DF)
  DF$xmax <- cumsum(DF$groupSum)
  DF$xmin <- DF$xmax - DF$groupSum
  DF$X <- row.names(DF)
  DF$groupSum <- NULL
  DF_melted <- melt(DF, id = c("X", "xmin", "xmax"), variable.name = "FILL")
  DF_melted <- DF_melted %>%
    group_by(X) %>%
    mutate(ymax = cumsum(value/sum(value)),
           ymin = ymax - value/sum(value))

  # Chi-sq test
  results <- chisq.test(table(data[[FILL]], data[[X]])) # fill and then x
  resid <- melt(results$residuals)
  names(resid) <- c("FILL", "X", "residual")


  # Merge data
  DF_all <- merge(DF_melted, resid)

  # Positions for labels
  DF_all$xposn <- DF_all$xmin + (DF_all$xmax - DF_all$xmin)/2
  index <- DF_all$xmax == max(DF_all$xmax)
  DF_all$yposn <- DF_all$ymin[index] + (DF_all$ymax[index] - DF_all$ymin[index])/2

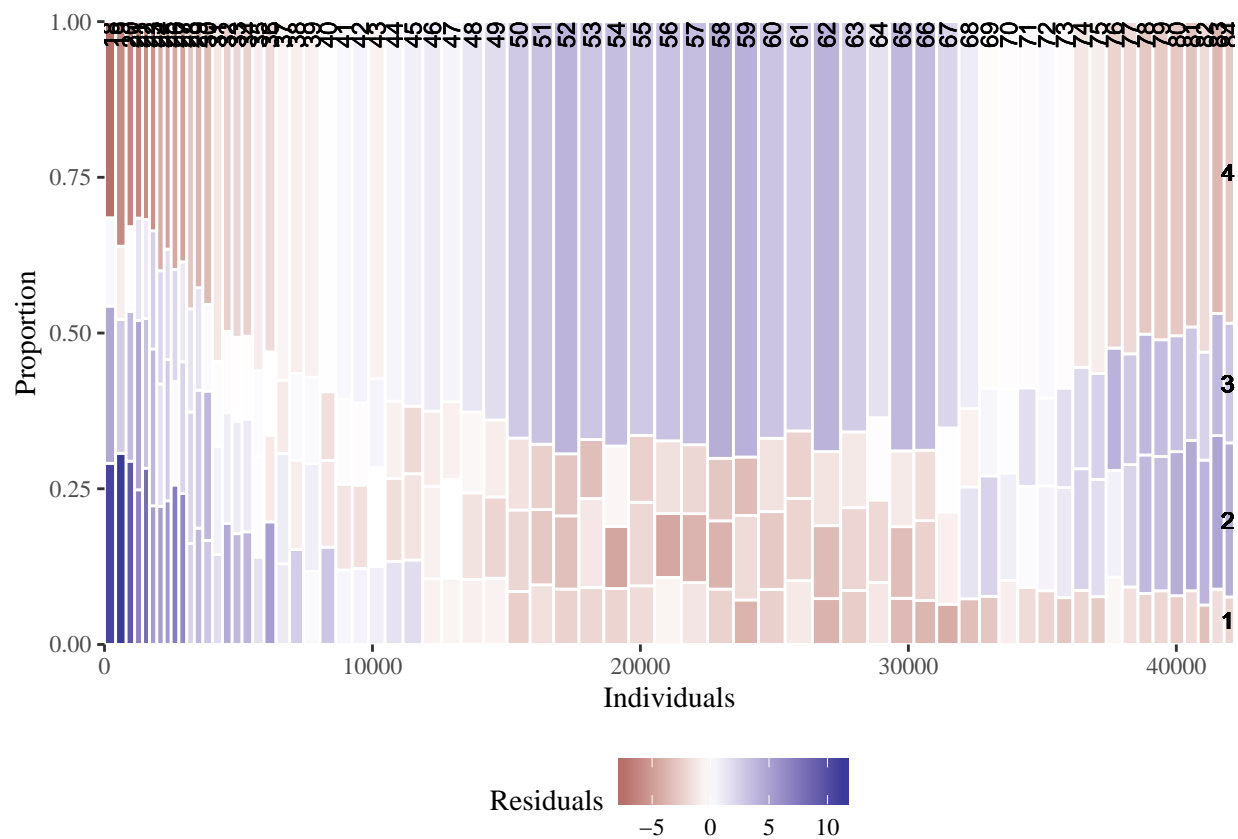
  # Plot
  g <- ggplot(DF_all, aes(ymin = ymin, ymax = ymax, xmin = xmin,
                        xmax = xmax, fill = residual)) +
    geom_rect(col = "white") +
    geom_text(aes(x = xposn, label = X),
              y = 1, size = 3, angle = 90, hjust = 1, show.legend = FALSE) +
    geom_text(aes(x = max(xmax), y = yposn, label = FILL),
              size = 3, hjust = 1, show.legend = FALSE) +
    scale_fill_gradient2("Residuals") +
    scale_x_continuous("Individuals", expand = c(0,0)) +
    scale_y_continuous("Proportion", expand = c(0,0)) +
    theme_tufte() +
    theme(legend.position = "bottom")
  print(g)
}

# BMI described by age (as previously seen)
mosaicGG(adult, X = "SRAGE_P", FILL = "RBMI")

```

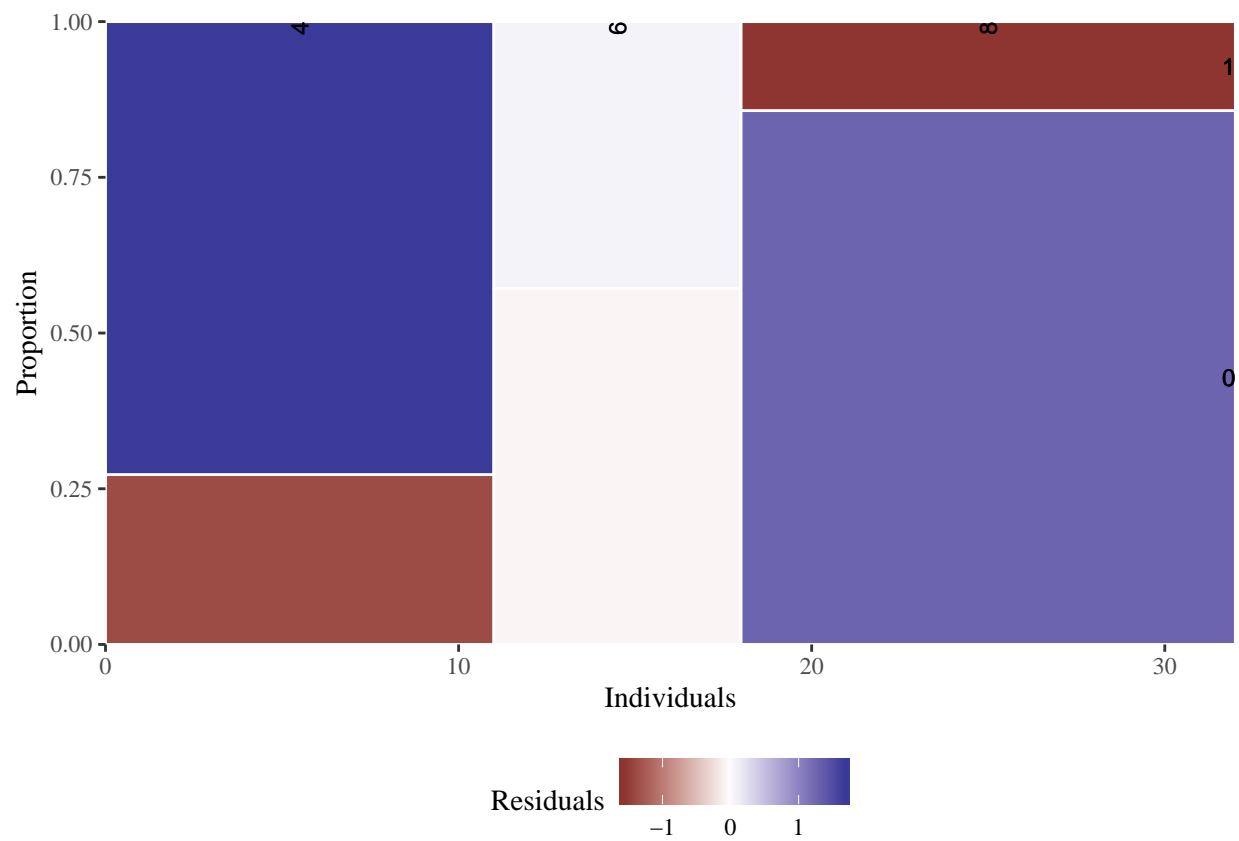


```
# Poverty described by age
mosaicGG(adult, X = "SRAGE_P", FILL = "POVLL")
```

```
# mtcars: am described by cyl
mosaicGG(mtcars, "cyl", "am")
```

```
## Warning in chisq.test(table(data[[FILL]], data[[X]]): Chi-squared
## approximation may be incorrect
```



```
# Vocab: vocabulary described by education
library(carData)
mosaicGG(Vocab, "education", "vocabulary")
```

```
## Warning in chisq.test(table(data[[FILL]], data[[X]])): Chi-squared
## approximation may be incorrect
```

